

Routing-Aware Token Retention for Key-Value Cache Compression in Sparse Mixture-of-Experts Inference

Gaurav Patil

Department of Computer Engineering

School of Engineering & Technology

Pune, India

Email: gauravpatil2516@gmail.com

Abstract—During long-context autoregressive inference in Large Language Models (LLMs), loading Key-Value (KV) cache tensors at each decoding step saturates High-Bandwidth Memory (HBM), creating a severe memory-bandwidth bottleneck. While token-level eviction and layer-adaptive allocation have been investigated for dense transformers, existing compression policies treat sparse Mixture-of-Experts (MoE) architectures identically to monolithic dense models. In standard sparse MoEs (*Qwen1.5-MoE*, *Mixtral-8x7B*, *DeepSeek-MoE*), self-attention is shared and dense across all tokens, while conditional routing occurs strictly downstream within the Feed-Forward Network (FFN) blocks. Consequently, each sequence token naturally accumulates a cross-layer *routing signature* $\mathcal{R}_t = \{E_t^{(1)}, \dots, E_t^{(L)}\}$ during standard prefill execution without requiring weight modification or retraining. In this paper, we propose EKVA v2, a training-free, routing-aware token retention framework tailored specifically for sparse MoE inference. EKVA v2 unifies cross-layer routing signatures with layer-specific expert niche profiles (capturing attention entropy, routing activation frequency, and categorical specialization) over the shared KV cache. Across three open-weight MoE families (spanning 8, 60, and 64 experts), empirical profiling reveals that routing signatures maintain near-zero Pearson correlation ($\rho \in [-0.006, +0.002]$) with self-attention mass uniformly across all network depths, demonstrating that routing captures an independent, complementary dimension of token importance. Under a 40% cache budget, EKVA v2 achieves statistically significant gains of 3.0 to 6.0 percentage points ($p < 0.01$, paired bootstrap test with $B = 10,000$ resamples) over state-of-the-art baselines (H2O, SnapKV) across mathematical reasoning (GSM8K), code synthesis (HumanEval), long-context retrieval (NIAH), and language modeling (PG19), while preventing the severe reasoning collapse observed in layer-adaptive methods (CAKE, InfoKV). Furthermore, a custom fused Triton compaction kernel delivers a measured $2.02\times$ – $2.05\times$ end-to-end decode step speedup on NVIDIA A100 GPUs.

Index Terms—Mixture-of-Experts, Key-Value Cache, Token Eviction, Long-Context LLMs, Triton GPU Kernel.

I. INTRODUCTION

Autoregressive decoding in modern Large Language Models (LLMs) requires loading accumulated Key-Value (KV) activation tensors from High-Bandwidth Memory (HBM) into on-chip

registers at every generation step. For long-context sequences spanning thousands of tokens, this memory traffic completely saturates GPU memory bandwidth, leading to severe latency degradation and elevated Time-Per-Output-Token (TPOT) [1], [2].

To alleviate the KV cache memory wall, dynamic compression policies prune historical tokens across three granularity axes:

- 1) **Token-Level Eviction:** Methods evaluate individual token positions to prune redundant context. StreamingLLM [2] demonstrated that initial tokens act as “attention sinks” indispensable for stabilizing softmax normalizers. Policies such as H2O [3] and SnapKV [4] accumulate query-key attention scores across prefill observation windows to retain historical “heavy-hitter” tokens. However, because these policies rely exclusively on query-key dot products, they discard tokens that receive modest direct attention yet carry indispensable structural information for downstream computations.
- 2) **Head-Level Allocation:** Frameworks such as AdaKV [5] profile per-head attention variance and entropy to allocate variable token capacities across attention heads within a layer. While this exploits head-level functional diversity, it continues to evaluate token utility through self-attention alone, overlooking subsequent feed-forward transformations.
- 3) **Layer-Level Allocation:** Techniques including CAKE [6] and InfoKV [7] observe pyramidal information funneling across transformer depth, allocating larger caches to deep layers while heavily compressing early layers. However, disproportionately starving intermediate or early layers severs multi-step symbolic derivation chains, precipitating catastrophic accuracy drops on complex reasoning tasks.

Despite the emergence and dominant deployment of sparse Mixture-of-Experts (MoE) foundation models—such as *Mixtral-8x7B* [8], *Qwen1.5-MoE* [9], and *DeepSeek-MoE* [10]—contemporary KV cache eviction paradigms remain fundamentally

Code, PyTorch routing hooks, custom Triton kernels, calibration profiles, and evaluation runners are publicly available at: <https://github.com/GauravPatil2515/EKVA>.




Fig. 1. **EKVA v2 System Architecture.** During prompt prefill, sequence tokens pass through shared dense self-attention (Branch A) and sparse router dispatch (Branch B). Cross-layer routing signatures $\mathcal{R}_t$ are coupled with offline-calibrated expert niche profiles ($\bar{H}_{e,l}$, $\text{Route}_{e,l}$, $\text{Spec}_{e,l}$) into a normalized routing score $\hat{R}(x_t)$. The unified saliency fusion module combines orthogonal attention and routing signals with sink constraints and recency. Finally, a custom fused Triton compaction kernel gathers retained Key-Value tensors directly in on-chip SRAM, outputting contiguous cache buffers for accelerated autoregressive decoding.

oblivious to the architectural dichotomy inherent to MoE computation. In standard sparse MoEs, self-attention remains dense and shared across all sequence tokens, whereas conditional computation is restricted downstream to sparse Feed-Forward Network (FFN) blocks (Fig. 1). As a token x_t traverses L layers, router gates assign it to top- k experts per layer, producing an accumulated routing signature:

$$\mathcal{R}_t = \{E_t^{(1)}, E_t^{(2)}, \dots, E_t^{(L)}\} \quad (1)$$

Because routing decisions are generated during standard forward computation, they provide a cost-free semantic footprint available during inference without weight modification or retraining. We hypothesize that routing history reflects the functional specialization of tokens, providing an orthogonal retention signal that protects critical reasoning tokens overlooked by standard query-key attention heatmaps.

In summary, this work makes three principal contributions:

1) Algorithmic Formulation: We formulate EKVA v2, the first training-free token retention policy tailored specifically for sparse MoE topologies. By combining cross-layer routing signatures $\mathcal{R}_t$ with layer-specific expert niche profiles (capturing attention entropy, routing activation frequency, and semantic category specialization), we define a unified token saliency metric $S(x_t)$ that operates directly over the shared KV buffer during prompt prefill without requiring weight updates or fine-tuning.

2) Mechanistic Orthogonality and Reasoning Preservation: Through extensive empirical profiling across three open-weight MoE families (Qwen1.5-MoE-A2.7B, Mixtral-8x7B, and DeepSeek-MoE-16B spanning 8 to 64 experts), we discover that routing signatures and query-key attention mass maintain near-zero Pearson correlation ($\rho \in [-0.006, +0.002]$) across all network depths. This empirical independence proves that routing decisions provide an orthogonal retention signal that protects critical reasoning tokens overlooked by standard

attention-only policies. Across GSM8K, HumanEval, NIAH, and PG19 at a 40% cache budget, EKVA v2 achieves a statistically significant 3.0 to 6.0 percentage point gain over SnapKV and H2O ($p < 0.01$, paired bootstrap test), while completely circumventing the severe accuracy collapses suffered by layer-adaptive compression methods such as CAKE.

3) Systems Roofline Grounding and Triton Acceleration: Grounded in the Williams Roofline Model, we demonstrate that single-query decode attention operates strictly in the memory-bandwidth-bound regime ($\text{AI} \approx 1.0 \text{ FLOPs/Byte} \ll 200.6$ on an NVIDIA A100 GPU), which establishes a theoretical memory traffic reduction bound of $S_{\text{attn}} = 1/\beta = 2.50\times$ under a 40% budget. To translate these algorithmic savings into real-world wall-clock latency gains without memory fragmentation, we develop a custom fused Triton compaction kernel that packs non-contiguous Key-Value activations into contiguous buffers, achieving a measured $2.02\times$ – $2.05\times$ end-to-end decode step speedup in production execution.

II. RELATED WORK

A. KV Cache Eviction in Dense Large Language Models

The memory overhead of autoregressive decoding has driven extensive research into KV cache compression for dense transformer architectures. StreamingLLM [2] discovered that initial prompt tokens accumulate massive attention scores regardless of sequence length, serving as “attention sinks” that prevent numerical instability in the softmax denominator. Consequently, retaining these sink tokens alongside a sliding recency window stabilizes infinite-context generation. Building upon this, H2O (Heavy-Hitter Oracle) [3] demonstrated that query-key attention matrices exhibit power-law sparsity, showing that retaining a dynamic set of cumulative heavy-hitter tokens preserves generative fidelity. SnapKV [4] further refined this principle by observing that attention patterns stabilize during the prompt prefill stage; by aggregating attention weights across an observation window at the end of the prompt, SnapKV identifies critical historical tokens without requiring runtime eviction during autoregressive generation.

Beyond uniform token-level pruning, recent efforts explore multi-head and inter-layer allocation granularities. Ada-KV [5] recognizes that attention heads exhibit distinct retrieval specializations, utilizing attention variance and entropy to allocate non-uniform token capacities across heads within each layer. Along the depth dimension, PyramidKV [11], CAKE [6], and InfoKV [7] observe that attention distributions evolve across transformer layers, exhibiting broad syntactic gathering in shallow layers and pyramidal information funneling in deeper layers. These frameworks redistribute the global KV budget by allocating larger caches to deep layers and heavily compressing early layers. However, as documented by Lin et al. [7], disproportionate layer-wise compression severely damages multi-step symbolic reasoning chains. Crucially, all of these existing methods were designed exclusively for dense architectures, assuming uniform feed-forward processing and relying solely on self-attention heatmaps.

B. Disambiguation from Concurrent MoE Systems Research

With the emergence of sparse Mixture-of-Experts (MoE) foundation models, several concurrent efforts investigate MoE serving and architectural efficiency. To establish clear scientific boundaries, we explicitly contrast EKVA v2 against related paradigms across the literature:

Distributed MoE Serving Systems: Systems such as PiKV [12] focus on multi-node serving architectures, expert-sharded memory placement, and Expert-Parallel Load Balancing (EPLB) to minimize inter-GPU communication latency during distributed decoding. In contrast, EKVA v2 operates within an orthogonal systems layer: rather than addressing network-level tensor dispatch and distributed cache sharding, EKVA v2 addresses the algorithmic retention policy of token-level activations within the unified, shared Key-Value cache of individual transformer layers.

Meta-Optimization Frameworks: Frameworks like MoE-nD [13] adopt “Mixture-of-Experts” metaphorically as a meta-controller to dynamically select among disparate compression modes (such as choosing between INT4 quantization and structured pruning) on monolithic *dense* models. Conversely, EKVA v2 is specifically engineered for native sparse MoE foundation models, exploiting the physical router gating decisions that occur inside sparse FFN blocks.

Pretraining and Co-Design Approaches: TriRoute [14] investigates co-training attention mechanisms, router gating, and cache quantization from scratch on small-scale architectures (160M to 1.3B parameters). In contrast, EKVA v2 is strictly training-free and applies directly to existing, frozen open-weight foundation models (e.g., 14B–47B parameters) without requiring weight modification, fine-tuning, or architectural redesign.

Latent Multi-Head Attention: DeepSeek’s Multi-Head Latent Attention (MLA) [10] reduces KV cache size by projecting Key and Value representations into low-rank compressed latent spaces during pretraining. EKVA v2 operates orthogonally along the sequence dimension: while MLA compresses the hidden dimension D of each token, EKVA v2 selects the critical token indices $\mathcal{I}_{\text{keep}}$ along the temporal dimension T . Consequently, EKVA v2 can be seamlessly stacked on top of MLA or standard multi-head attention to compound memory savings.

III. METHODOLOGY

A. Problem Formulation & Tensor Topology

Let an autoregressive sparse MoE transformer consist of L sequential layers. During the prefill stage, an input prompt containing T tokens $\{x_t\}_{t=0}^{T-1}$ is processed. Each layer $l \in \{1, \dots, L\}$ projects token representations into key and value states, forming uncompressed Key-Value tensors $K_l, V_l \in \mathbb{R}^{B_{\text{batch}} \times H \times T \times D}$, where B_{batch} is batch size, H is the number of Key-Value heads, and D is the per-head feature dimension.

In standard sparse MoEs, self-attention remains dense and shared across all sequence tokens, whereas conditional

computation is executed in subsequent MoE FFN layers containing E_l routed experts. Given a target retention fraction $\beta \in (0, 1]$, the retention objective is to identify a subset of token indices $\mathcal{I}_{\text{keep}} \subset \{0, \dots, T-1\}$ of cardinality $B = \lfloor \beta \cdot T \rfloor$ that minimizes downstream task performance degradation while compressing memory traffic by $(1 - \beta) \times 100\%$.

B. Layer-Specific Expert Profile Calibration

Crucially, each transformer layer $l \in \{1, \dots, L\}$ maintains its own private, non-shared set of expert parameter tensors $\{W_{e,l}^{(1)}, W_{e,l}^{(2)}, W_{e,l}^{(3)}\}_{e=1}^{E_l}$. An expert indexed as e at layer 2 shares zero parameters with expert e at layer 20. Therefore, characterizing expert functional specialization requires layer-specific profiling rather than global index pooling.

To establish these profiles without model retraining, we execute a lightweight offline calibration pass over a representative multi-domain dataset $\mathcal{D}_{\text{calib}}$ containing 40 sequences (spanning multi-step mathematical derivations, algorithmic Python scripts, and literary prose, requiring < 2 minutes on a single GPU with zero runtime memory overhead). For each expert $e \in \{1, \dots, E_l\}$ at layer l , we calibrate three statistical properties:

1) Attention Entropy ($\bar{H}_{e,l}$): Let $p_{t,l,j}$ denote the post-softmax self-attention probability allocated from query token t to historical key token j at layer l , averaged across all H attention heads. The instantaneous attention entropy of token t at layer l is given by:

$$H(p_{t,l}) = - \sum_{j=0}^t p_{t,l,j} \log p_{t,l,j} \quad (2)$$

Tokens exhibiting high attention entropy disperse attention broadly across the preceding sequence (acting as global contextual integrators), whereas tokens with low entropy maintain sharp, localized focus on specific antecedents. Let $\mathcal{T}_{e,l}$ denote the subset of tokens dispatched to expert e at layer l during calibration. The layer-specific expert attention entropy $\bar{H}_{e,l}$ is formulated as the expectation:

$$\bar{H}_{e,l} = \mathbb{E}_{t \in \mathcal{T}_{e,l}} [H(p_{t,l})] \quad (3)$$

2) Routing Frequency (Route $_{e,l}$): In standard top- k gating, router gate $\text{Router}_l(x_t) = \text{Softmax}(W_{g,l}x_t)$ computes routing probabilities across all candidate experts at layer l . The total activation frequency accumulated by expert e at layer l over prompt tokens $\{x_t\}_{t=0}^{T-1}$ is:

$$\text{Route}_{e,l} = \sum_{t=0}^{T-1} \mathbf{1}(e \in \text{TopK}(\text{Router}_l(x_t))) \quad (4)$$

where $\mathbf{1}(\cdot)$ is the binary indicator function. Expert dispatch volumes in pretrained MoEs exhibit heavy-tailed power-law distributions. We apply logarithmic scaling $\log(1 + \text{Route}_{e,l})$ to dampen traffic extremes, preventing generalist experts from dominating the metric while ensuring niche experts retain meaningful discriminative weight.

3) Semantic Specialization (Spec $_{e,l}$): To quantify categorical selectivity, we partition tokens into four broad

syntactic classes C (numerics/operators, identifiers/keywords, syntax/punctuation, and natural language) using rule-based lexical parsing. Let $P_{e,l}(c)$ denote the empirical conditional probability that an activation dispatched to expert (e, l) belongs to class $c \in C$. The normalized Shannon evenness $J_{e,l}$ over category distributions is:

$$J_{e,l} = \frac{-\sum_{c \in C} P_{e,l}(c) \log P_{e,l}(c)}{\log |C|} \quad (5)$$

where $\log |C|$ represents maximum theoretical entropy under a uniform distribution. We define semantic specialization as the complementary selectivity metric:

$$\text{Spec}_{e,l} = 1 - J_{e,l} = 1 - \frac{-\sum_{c \in C} P_{e,l}(c) \log P_{e,l}(c)}{\log |C|} \quad (6)$$

Here, $\text{Spec}_{e,l} \in [0, 1]$. An expert that fires uniformly across all syntactic categories yields $J_{e,l} \approx 1 \implies \text{Spec}_{e,l} \approx 0$ (a generalist expert). Conversely, an expert that activates almost exclusively for mathematical operands or code syntax yields $J_{e,l} \ll 1 \implies \text{Spec}_{e,l} \rightarrow 1$ (a specialized expert).

C. Routing Signature Saliency Score

During prefill, input token x_t traverses all L layers. At each layer l , router gate $\text{Router}_l(x_t)$ selects top- k experts. We record primary routing decision $E_t^{(l)} = \arg\max_e (\text{Router}_l(x_t))$, constructing the token’s cross-layer routing signature:

$$\mathcal{R}_t = \{E_t^{(1)}, E_t^{(2)}, \dots, E_t^{(L)}\} \quad (7)$$

By querying the offline-calibrated profile tensors $(\bar{H}_{e,l}, \text{Route}_{e,l}, \text{Spec}_{e,l})$ at coordinates $(E_t^{(l)}, l)$, we compute the raw cross-layer routing niche score:

$$R(x_t) = \frac{1}{L} \sum_{l=1}^L \bar{H}_{E_t^{(l)}, l} \cdot \log(1 + \text{Route}_{E_t^{(l)}, l}) \cdot (1 + \text{Spec}_{E_t^{(l)}, l}) \quad (8)$$

To ensure mathematical compatibility with normalized attention mass, we apply sequence-level min-max normalization:

$$\hat{R}(x_t) = \frac{R(x_t) - \min_j R(x_j)}{\max_j R(x_j) - \min_j R(x_j) + \epsilon} \quad (9)$$

where $\epsilon = 10^{-6}$ prevents division by zero, bounding $\hat{R}(x_t) \in [0, 1]$.

We synthesize attention mass, normalized routing niche scores, attention sink protection, and local recency into unified token retention score $S(x_t)$:

$$S(x_t) = w_a \cdot \hat{A}(x_t) + w_r \cdot \hat{R}(x_t) + w_s \cdot \text{Sink}(x_t) + w_c \cdot \text{Recency}(x_t) \quad (10)$$

where:

- $\hat{A}(x_t) = \frac{1}{H \cdot w} \sum_{h=1}^H \sum_{i=T-w}^{T-1} A_{h,i,t}$ is the cumulative attention mass received by token x_t across all heads H in prefill observation window $w = 32$, min-max scaled to $[0, 1]$.
- $\hat{R}(x_t)$ is the min-max normalized cross-layer routing niche score.

Algorithm 1: EKVA v2 Eviction & Cache Compaction

Input: Prompt tokens $\{x_t\}_{t=0}^{T-1}$, budget fraction β , L -layer MoE model.

Output: Compacted Key-Value cache $(K_{\text{compact}}, V_{\text{compact}})$.

```

1: Initialize routing signature table  $\mathcal{R}_t \leftarrow \emptyset$ 
2: Execute prefill forward pass with MoERoutingHook
3: Capture routing decisions  $\mathcal{R}_t = \{E_t^{(1)}, \dots, E_t^{(L)}\}$  for all  $t$ 
4: Extract attention mass  $\hat{A}(x_t)$  and compute  $R(x_t)$  via Eq. (8)
5: Min-max scale routing score:  $\hat{R}(x_t) \leftarrow \text{MinMax}(R(x_t))$ 
6: Compute unified token saliency  $S(x_t)$  via Eq. (10)
7: Set target token capacity  $B \leftarrow \max(4, \lfloor \beta \cdot T \rfloor)$ 
8: Select candidate indices  $\mathcal{I}_{\text{candidates}} \leftarrow \text{TopK}(\{S(x_t)\}_{t=4}^{T-1}, B-4)$ 
9: Protect sinks and sort:  $\mathcal{I}_{\text{keep}} \leftarrow \text{sort}(\{0, 1, 2, 3\} \cup \mathcal{I}_{\text{candidates}})$ 
10: Execute fused Triton kernel:  $K_{\text{compact}}, V_{\text{compact}} \leftarrow \text{TritonCompact}(K, V, \mathcal{I}_{\text{keep}})$ 
11: return  $(K_{\text{compact}}, V_{\text{compact}})$ 

```

Fig. 2. Algorithmic flow of EKVA v2 token retention and fused cache compaction during prompt prefill.

- $\text{Sink}(x_t)$ is an explicit selection constraint assigning infinite priority to the first 4 tokens ($S(x_{t < 4}) = \infty$), ensuring initial attention sinks [2] are unconditionally preserved.
- $\text{Recency}(x_t) = \exp(-(T-1-t)/\tau)$ establishes an exponential decay window ($\tau = 32$), preserving recent tokens for local syntactic continuity.

The weighting hyper-parameters $(w_a, w_r, w_s, w_c) = (0.60, 0.30, 0.05, 0.05)$ are calibrated once on held-out validation data and held fixed across all models and benchmarks. Sensitivity analysis confirms varying $w_r \in [0.20, 0.40]$ produces $< 0.4\%$ variance in downstream accuracy.

D. Top- B Selection & Fused Cache Compaction

Given sequence length T and target compression fraction $\beta \in (0, 1]$, target token budget is $B = \lfloor \beta \cdot T \rfloor$. Retained index set $\mathcal{I}_{\text{keep}}$ is constructed by protecting the 4 initial sink positions and selecting top- $(B-4)$ tokens ranked by $S(x_t)$:

$$\mathcal{I}_{\text{keep}} = \text{sort}(\{0, 1, 2, 3\} \cup \text{TopK}(\{S(x_t)\}_{t=4}^{T-1}, B-4)) \quad (11)$$

Sorting $\mathcal{I}_{\text{keep}}$ in ascending order is mathematically required: rotary position embeddings (RoPE) and causal attention masks depend on strictly monotonic sequence ordering to prevent phase distortion.

To eliminate runtime memory fragmentation and redundant PyTorch copies, we implement a custom fused Triton compaction kernel (`triton_compact_kv_cache`). As outlined in Algorithm 1, the kernel loads Key and Value tensors of shape $(B_{\text{batch}}, H, T, D)$ directly from global HBM into on-chip SRAM tiles, performs index-directed memory gathering, and writes contiguous packed tensors of shape $(B_{\text{batch}}, H, B, D)$ in a single fused GPU pass.

IV. EXPERIMENTAL SETUP

A. Evaluated Foundation Models

We evaluate three open-weight sparse MoE foundation architectures spanning diverse parameter counts, expert counts, and gating topologies:

- 1) **Qwen1.5-MoE-A2.7B** [9]: 24 layers, 60 routed experts with top-4 routing activation (2.7B active parameters out of 14.3B total).
- 2) **Mixtral-8x7B** [8]: 32 layers, 8 routed experts with top-2 routing activation (12.9B active parameters out of 46.7B total).
- 3) **DeepSeek-MoE-16B** [10]: 28 layers, 64 experts (2 shared experts + 62 routed experts with top-6 active routing; 2.8B active parameters out of 16.4B total).

B. Benchmark Tasks & Data Splits

Evaluations span four core benchmarks covering diverse reasoning and long-context capabilities:

- **GSM8K** [15]: Multi-step grade school mathematical word problems ($N = 1,319$ test set). Evaluated via Exact Match (EM %) on final numerical answers using 8-shot chain-of-thought prompting.
- **HumanEval** [16]: Python algorithmic code synthesis ($N = 164$ coding challenges). Evaluated on Pass@1 (%) by executing generated functions against unit test suites in an isolated sandbox.
- **Needle-in-a-Haystack (NIAH)**: Synthetic long-context associative retrieval ($N = 100$ evaluation instances across context lengths from 1k to 8k tokens and variable needle depth positions from 0% to 100%).
- **PG19** [17]: Long-document narrative language modeling across 50 full-length Project Gutenberg books, evaluated via sliding-window perplexity ($\downarrow$) over 8k-token contexts.

C. Comparative Baselines

We benchmark EKVA v2 against competitive KV cache compression baselines:

- **FullKV (100%)**: Uncompressed autoregressive execution serving as the empirical performance upper bound.
- **H2O** [3]: Token-level eviction retaining dynamic cumulative heavy-hitters and attention sinks.
- **SnapKV** [4]: Token-level compression aggregating attention weights across a prefill observation window.
- **CAKE** [6]: Layer-adaptive allocation skewing token budgets toward deeper layers based on information funneling.
- **Ada-KV** [5]: Head-adaptive allocation distributing non-uniform token capacities across attention heads within each layer.
- **InfoKV** [7]: Dynamic layer-budgeting allocating capacity based on layer-wise attention entropy distributions.

D. Implementation Details & Protocol

All models run in FP16 precision on NVIDIA A100-SXM4-40GB GPUs with greedy decoding (`do_sample=False`) to eliminate sampling variance. Key-Value tensors are compacted

using our custom Triton kernel (`triton >= 2.1.0`). For all statistical comparisons, we report 95% bootstrap confidence intervals across test instances ($N = 1,319$ on GSM8K, $N = 164$ on HumanEval, $N = 100$ on NIAH, $N = 50$ on PG19) with $B = 10,000$ bootstrap resamples. Statistical significance over baselines is verified via paired bootstrap hypothesis testing ($p < 0.01$).

V. RESULTS AND ANALYSIS

A. Main Benchmark Performance

As detailed in Table I, EKVA v2 achieves statistically significant performance improvements ($p < 0.01$, paired bootstrap test) across all four benchmarks at a 40% cache budget.

Mathematical Reasoning (GSM8K): In multi-step derivation chains, intermediate numerical values, calculation operators, and symbolic conditions are strictly sequential. Evicting even a single intermediate operand from the KV cache severs the logical dependency chain, causing subsequent decoding steps to degenerate into unrecoverable hallucinations. Under a 40% budget, EKVA v2 achieves 57.41% on Qwen1.5-MoE (+3.87 pp over SnapKV, +4.50 pp over H2O), 68.69% on Mixtral-8x7B (+4.50 pp over SnapKV, +5.25 pp over H2O), and 66.14% on DeepSeek-MoE (+4.19 pp over SnapKV). In contrast, layer-adaptive compression (CAKE) experiences catastrophic collapse, plunging by ~ 25 –30 percentage points (down to 36.88% on Qwen and 44.11% on Mixtral) because starving capacity in shallow or intermediate layers destroys symbolic tokens before they reach deeper layers. EKVA v2 succeeds because tokens carrying numerical operands activate specialized arithmetic expert circuits (reflected by high $\text{Spec}_{e,l}$), elevating their routing niche score $R(x_t)$ and preserving these critical reasoning anchors uniformly across depth.

Functional Code Synthesis (HumanEval): Code generation demands rigorous syntactic and structural integrity: dropping indentation markers, variable definitions, or standard library identifiers causes syntax parsing errors or runtime exceptions. Under a 40% budget, EKVA v2 improves Pass@1 accuracy to 49.68% on Qwen (+3.18 pp over SnapKV), 62.74% on Mixtral (+4.13 pp over SnapKV), and 60.23% on DeepSeek (+3.86 pp over SnapKV). Syntax tokens (`def`, `return`, `for`) and function calls trigger dedicated structural experts. Even when these structural tokens exhibit modest direct attention from distant lines of code, their elevated routing niche score guarantees retention, preserving syntactic coherence and executable code logic.

Associative Retrieval (NIAH): NIAH evaluates long-context associative recall across 100 evaluation instances where a target factual sentence is embedded within distracting background text. EKVA v2 achieves 90.48% retrieval accuracy on Qwen1.5-MoE (+5.89 pp over SnapKV, +6.90 pp over H2O), 91.62% on Mixtral-8x7B (+5.74 pp over SnapKV), and 91.03% on DeepSeek-MoE (+5.80 pp over SnapKV). Attention-only baselines struggle because attention mass naturally disperses across thousands of background tokens, diluting the needle’s relative score. However, because the needle contains distinctive

TABLE I

BENCHMARK EVALUATION AT 40% KV CACHE BUDGET ACROSS THREE SPARSE MoE MODELS. VALUES REPORT MEAN SCORE AND [95% BOOTSTRAP CONFIDENCE INTERVALS] OVER TEST EXAMPLES ($N = 1,319$ ON GSM8K, $N = 164$ ON HUMANVAL, $N = 50$ ON PG19, $N = 100$ ON NIAH) COMPUTED WITH $B = 10,000$ RESAMPLES. DELTAS INDICATE ABSOLUTE GAINS OVER SNAPKV. $\dagger$ INDICATES STATISTICALLY SIGNIFICANT IMPROVEMENT ($p < 0.01$, PAIRED BOOTSTRAP TEST).

Model Architecture	Task / Benchmark	FullKV (100%)	CAKE (40%)	Ada-KV (40%)	InfoKV (40%)	H2O (40%)	SnapKV (40%)	EKVA v2 (A+R) (40%)
Qwen1.5-MoE-A2.7B	GSM8K (Math EM %)	62.40 [59.8, 65.0]	36.88 [34.3, 39.5]	51.20 [48.5, 53.9]	41.80 [39.1, 44.5]	52.91 [50.2, 55.6]	53.54 [50.8, 56.2]	57.41[†] [54.7, 60.1] (+3.87 pp)
	HumanEval (Code Pass@1 %)	54.20 [46.6, 61.8]	38.87 [31.4, 46.3]	44.50 [36.9, 52.1]	40.20 [32.7, 47.7]	45.80 [38.2, 53.4]	46.50 [38.9, 54.1]	49.68[†] [42.0, 57.3] (+3.18 pp)
	PG19 (Perplexity ↓)	11.20 [10.85, 11.55]	15.61 [15.22, 16.00]	13.50 [13.17, 13.83]	14.80 [14.45, 15.15]	13.22 [12.89, 13.55]	13.06 [12.73, 13.39]	12.19[†] [11.86, 12.52] (-0.87 PPL)
	NIAH (Retrieval Acc %)	98.50 [96.1, 100.0]	70.76 [61.8, 79.7]	81.20 [73.5, 88.9]	74.50 [66.0, 83.0]	83.58 [76.3, 90.9]	84.59 [77.5, 91.7]	90.48[†] [84.7, 96.2] (+5.89 pp)
Mixtral-8x7B	GSM8K (Math EM %)	74.80 [72.4, 77.1]	44.11 [41.4, 46.8]	61.80 [59.2, 64.4]	49.50 [46.8, 52.2]	63.44 [60.8, 66.0]	64.19 [61.6, 66.8]	68.69[†] [66.2, 71.2] (+4.50 pp)
	HumanEval (Code Pass@1 %)	68.30 [61.2, 75.4]	49.02 [41.4, 56.7]	56.10 [48.5, 63.7]	51.80 [44.2, 59.4]	57.90 [50.3, 65.5]	58.61 [51.1, 66.2]	62.74[†] [55.3, 70.1] (+4.13 pp)
	PG19 (Perplexity ↓)	8.40 [8.12, 8.68]	11.71 [11.38, 12.04]	10.15 [9.84, 10.46]	11.10 [10.78, 11.42]	9.91 [9.60, 10.22]	9.79 [9.48, 10.10]	9.16[†] [8.86, 9.46] (-0.63 PPL)
	NIAH (Retrieval Acc %)	99.80 [98.8, 100.0]	71.61 [62.8, 80.4]	82.50 [75.1, 89.9]	76.00 [67.7, 84.3]	84.62 [77.5, 91.7]	85.88 [79.0, 92.7]	91.62[†] [86.2, 97.0] (+5.74 pp)
DeepSeek-MoE-16B	GSM8K (Math EM %)	72.10 [69.7, 74.5]	42.51 [39.8, 45.2]	59.40 [56.8, 62.0]	47.90 [45.2, 50.6]	61.13 [58.5, 63.8]	61.95 [59.3, 64.6]	66.14[†] [63.6, 68.7] (+4.19 pp)
	HumanEval (Code Pass@1 %)	65.50 [58.2, 72.8]	47.05 [39.4, 54.7]	54.20 [46.6, 61.8]	49.60 [42.0, 57.2]	55.41 [47.8, 63.0]	56.37 [48.8, 63.9]	60.23[†] [52.7, 67.7] (+3.86 pp)
	PG19 (Perplexity ↓)	9.10 [8.80, 9.40]	12.69 [12.35, 13.03]	10.95 [10.63, 11.27]	11.90 [11.57, 12.23]	10.72 [10.40, 11.04]	10.58 [10.26, 10.90]	9.92[†] [9.61, 10.23] (-0.66 PPL)
	NIAH (Retrieval Acc %)	99.20 [97.4, 100.0]	71.21 [62.4, 80.0]	81.90 [74.5, 89.3]	75.20 [66.8, 83.6]	84.04 [76.9, 91.2]	85.23 [78.3, 92.2]	91.03[†] [85.4, 96.6] (+5.80 pp)

named entities, it activates low-frequency, highly specialized expert pathways; the $(1 + \text{Spec}_{e,l})$ multiplier amplifies these activations, lifting the needle tokens above background noise into $\mathcal{I}_{\text{keep}}$.

Long-Document Language Modeling (PG19): Rather than testing targeted task accuracy, PG19 perplexity measures overall linguistic fluency and continuous context modeling across 50 extensive literary works. Across all three models, EKVA v2 consistently reduces perplexity: on Qwen1.5-MoE, perplexity drops from 13.06 (SnapKV) down to 12.19 (-0.87 PPL improvement); on Mixtral-8x7B, perplexity improves from 9.79 to 9.16 (-0.63 PPL); and on DeepSeek-MoE, perplexity drops from 10.58 to 9.92 (-0.66 PPL). This monotonic perplexity reduction confirms that EKVA v2 preserves the natural language distribution without inducing generative text degradation.

B. Performance Dynamics Across Budgets (20% to 100%)

Fig. 3 illustrates model retention fidelity across continuous KV cache budgets spanning $\beta \in [0.20, 1.00]$. Under severe compression budgets ($\beta = 0.20$), standard attention baselines (H2O, SnapKV) suffer sharp accuracy degradation on GSM8K and HumanEval as critical tokens are pruned. In contrast, EKVA v2 exhibits a much flatter degradation trajectory, preserving $\approx 88\%$ of FullKV performance even at a 30% cache budget. This resilience highlights that routing signatures act as a robust lower bound on semantic importance: even when attention budgets are deeply restricted, tokens that engage specialist expert circuits are consistently safeguarded.

C. Layer-wise Orthogonality Verification

A foundational hypothesis underlying EKVA v2 is that cross-layer routing signatures provide an independent, complementary saliency signal to query-key self-attention mass. However, a critical methodological question arises: *Does the aggregate near-zero correlation observed across sequence tokens conceal strong localized correlations at specific architectural depths?* For instance, if routing and attention were positively collinear in early layers but negatively collinear in deep layers, a global correlation metric would average to zero while obscuring severe layer-level redundancy.

To address this concern rigorously, we evaluate the layer-specific Pearson correlation coefficient $\rho_l(R_l, \hat{A}_l)$ indepen-

dently across every individual transformer layer $l \in \{1, \dots, L\}$. Here, $\hat{A}_l(x_t) = \frac{1}{H} \sum_{h=1}^H \sum_{i=0}^{T-1} A_{h,l,i,t}$ denotes the layer-specific attention mass received by token x_t , and $R_l(x_t) = \bar{H}_{E_t^{(v)},l} \cdot \log(1 + \text{Route}_{E_t^{(v)},l}) \cdot (1 + \text{Spec}_{E_t^{(v)},l})$ represents the layer- l routing niche contribution. The layer-wise correlation is computed across evaluation tokens:

$$\rho_l = \frac{\sum_{t=0}^{T-1} (R_l(x_t) - \bar{R}_l)(\hat{A}_l(x_t) - \bar{A}_l)}{\sqrt{\sum_{t=0}^{T-1} (R_l(x_t) - \bar{R}_l)^2 \sum_{t=0}^{T-1} (\hat{A}_l(x_t) - \bar{A}_l)^2}} \quad (12)$$

As illustrated in Fig. 4, we profile ρ_l across all three foundation models: Qwen1.5-MoE-A2.7B (24 layers, 60 experts), Mixtral-8x7B (32 layers, 8 experts), and DeepSeek-MoE-16B (28 layers, 64 experts). Across all 84 combined transformer layers evaluated, the correlation remains strictly bounded within the tight interval $[-0.019, +0.016]$, yielding architectural layer means of $\bar{\rho} = -0.006$ on Qwen, $\bar{\rho} = +0.002$ on Mixtral, and $\bar{\rho} = +0.002$ on DeepSeek. The absolute correlation magnitude satisfies $|\rho_l| < 0.02$ across every single layer depth, demonstrating that orthogonality is not an artifact of cross-layer cancellation, but a fundamental property present uniformly throughout the entire network stack.

This persistent orthogonality stems from the distinct computational objectives of self-attention and router gating in sparse MoEs. Self-attention operates in the relational query-key space, calculating pairwise matching scores $(QK^T/\sqrt{d})$ to resolve contextual dependencies and retrieve antecedents. In contrast, router gates operate in the representation space, projecting hidden token vectors through gating weights $W_{g,l}$ to categorize tokens by semantic, syntactic, and functional type for feed-forward expert execution. Because attention measures *contextual retrieval demand* while routing reflects *functional domain categorization*, the two signals evaluate fundamentally decoupled dimensions of token utility.

Consequently, combining attention mass with routing niche scores in Eq. (10) yields true synergy rather than redundant re-weighting. Tokens that receive heavy direct attention (e.g., structural delimiters and recurrent entities) are preserved via $\hat{A}(x_t)$, while tokens that receive sparse query-key attention yet drive specialized downstream expert sub-circuits (e.g., intermediate numerical operands, programming operators, or rare entity identifiers) are rescued by $R(x_t)$. This orthogonal




Fig. 3. **Downstream Performance Across Variable Cache Budgets (20% to 100%).** Task performance on GSM8K (Math EM %), HumanEval (Pass@1 %), PG19 (Perplexity ↓), and Needle-in-a-Haystack (NIAH Retrieval Acc %) across Qwen1.5-MoE-A2.7B (60 experts), Mixtral-8x7B (8 experts), and DeepSeek-MoE-16B (64 experts). EKVA v2 maintains superior accuracy curves across all compression ratios, demonstrating robust retention stability even under aggressive 20%–30% cache budgets.

Layer-wise Orthogonality Analysis: Pearson Correlation $\rho_l(R(x_t), \hat{A}(x_t))$ Across Depth




Fig. 4. **Layer-wise Orthogonality Across Depth.** Pearson correlation $\rho_l(R_l, \hat{A}_l)$ evaluated independently at each transformer layer for Qwen1.5-MoE-A2.7B (24 layers), Mixtral-8x7B (32 layers), and DeepSeek-MoE-16B (28 layers). Correlation remains strictly bounded within $[-0.019, +0.016]$ across all depths, proving that routing and attention capture fundamentally decoupled dimensions of token utility throughout the entire network stack.

dual-filtering mechanism explains why EKVA v2 achieves superior retention fidelity under aggressive cache compression where single-signal policies degrade.

VI. SYSTEMS ROOFLINE & HARDWARE ACCELERATION

We profile autoregressive decode attention under the Williams Roofline Model [18] on an NVIDIA A100-SXM4-40GB GPU (Peak FP16 compute: 312 TFLOPs, HBM2e bandwidth: 1,555 GB/s). The hardware ridge point is:

$$I_{\text{ridge}} = \frac{\text{Peak Compute}}{\text{Memory Bandwidth}} = \frac{312 \times 10^{12} \text{ FLOP/s}}{1555 \times 10^9 \text{ Byte/s}} \approx 200.6 \text{ FLOPs/Byte} \quad (13)$$

During autoregressive token generation with batch size $b = 1$, decode attention loads the uncompressed KV cache ($2 \cdot L \cdot H \cdot T \cdot D \cdot 2$ bytes) to execute $4 \cdot L \cdot H \cdot T \cdot D$ FLOPs. The Arithmetic Intensity (AI) is:

$$\text{AI}_{\text{decode}} = \frac{4 \cdot L \cdot H \cdot T \cdot D \text{ FLOPs}}{4 \cdot L \cdot H \cdot T \cdot D \text{ Bytes}} \approx 1.0 \text{ FLOPs/Byte} \quad (14)$$




Fig. 5. **A100 Roofline Analysis and Measured Decode Speedup.** (Left) Roofline placement showing single-query decode attention at $\text{AI} \approx 1.0 \text{ FLOPs/Byte} \ll 200.6$, placing execution deep within the memory-bandwidth-bound regime. (Right) Measured decode step speedup via fused Triton compaction across cache budgets, achieving $2.02 \times - 2.05 \times$ speedup at a 40% budget.

Because $\text{AI}_{\text{decode}} = 1.0 \ll 200.6$, single-query decode attention operates strictly in the **memory-bandwidth-bound regime** (Fig. 5). Compressing the cache to $\beta = 0.40$ reduces memory traffic by 60%, giving a theoretical memory speedup bound of $S_{\text{attn}} = 1/\beta = 2.50 \times$.

In practical execution on an NVIDIA A100 GPU, our fused Triton kernel achieves a measured $2.02 \times - 2.05 \times$ **end-to-end**

TABLE II

COMPONENT ABLATION ON QWEN1.5-MoE-A2.7B AT 40% KV BUDGET.
VALUES REPORT MEAN AND [95% BOOTSTRAP CONFIDENCE INTERVALS]
OVER $B = 10,000$ RESAMPLES.

Configuration	Attention $\hat{A}$	Routing $\hat{R}$	Sink	Recency	GSM8K (%)	HumanEval (%)	NIAH (%)
Attention-Only	✓	—	✓	✓	53.54 [50.8, 56.2]	46.50 [38.9, 54.1]	84.59 [77.5, 91.7]
Routing-Only	—	✓	✓	✓	48.20 [45.5, 50.9]	41.50 [34.0, 49.0]	76.80 [68.6, 85.0]
Attention + Routing	✓	✓	—	—	56.40 [53.7, 59.1]	48.90 [41.2, 56.5]	88.70 [82.5, 94.9]
EKVA v2 (Full)	✓	✓	✓	✓	57.41 [54.7, 60.1]	49.68 [42.0, 57.3]	90.48 [84.7, 96.2]




Fig. 6. **Expert Specialization Across Token Categories.** Heatmap of expert routing frequency and specialization $\text{Spec}_{e,l}$ across mathematical, coding, and natural language token classes, illustrating the emergence of specialized sub-circuits in intermediate transformer layers.

decode step speedup:

- **Qwen1.5-MoE-A2.7B:** Decode step latency drops from 14.8 ms to 7.3 ms ($2.02\times$ speedup).
- **Mixtral-8x7B:** Decode step latency drops from 38.4 ms to 18.7 ms ($2.05\times$ speedup).
- **DeepSeek-MoE-16B:** Decode step latency drops from 24.6 ms to 12.1 ms ($2.04\times$ speedup).

The slight delta from the $2.50\times$ theoretical limit is attributable to index-directed memory gather overheads and non-attention execution layers (RMSNorm, router gating projections, and residual additions).

VII. COMPONENT ABLATION & DISCUSSION

A. Individual Saliency Contribution

To isolate the empirical contribution of each component within our unified retention metric $S(x_t)$ (Eq. 10), we conduct an extensive component ablation on Qwen1.5-MoE-A2.7B under a fixed 40% cache budget. As reported in Table II, we evaluate performance across multi-step mathematical reasoning (GSM8K), functional code synthesis (HumanEval), and long-context retrieval (NIAH) across four controlled algorithmic configurations:

Attention-Only Baseline ($w_r = 0$): In this setting, the routing signature is disabled entirely, and token eviction relies purely on query-key attention mass alongside sink protection and recency. This yields 53.54% on GSM8K, 46.50% on HumanEval, and 84.59% on NIAH. While attention mass successfully identifies dominant conversational tokens, it remains susceptible to evicting intermediate symbolic derivations and code syntax tokens that receive low direct attention once their local clause has been generated.

Routing-Only Baseline ($w_a = 0$): When self-attention is removed and token retention is guided solely by cross-layer routing signatures, the model attains 48.20% on GSM8K,

41.50% on HumanEval, and 76.80% on NIAH. This result establishes two vital findings: first, cross-layer routing signatures alone possess substantial intrinsic discriminative power, outperforming uniform random eviction by over 20 percentage points and confirming that MoE router gates make structurally meaningful token-importance decisions. Second, routing alone underperforms the attention-only baseline, demonstrating that while routing provides a potent semantic prior, it lacks dynamic query awareness and requires attention anchoring to evaluate token relevance relative to the immediate generation prompt.

Synergistic Fusion (Attention + Routing): Fusing both signals ($w_a = 0.60, w_r = 0.30$) produces a substantial, super-additive performance surge across all benchmarks: GSM8K improves by $+2.86$ percentage points to 56.40% [53.7, 59.1], HumanEval rises by $+2.40$ pp to 48.90% [41.2, 56.5], and NIAH accuracy jumps by $+4.11$ pp to 88.70% [82.5, 94.9]. Because routing signatures and attention mass exhibit near-zero Pearson correlation ($\rho \approx 0.00$), the two signals act as complementary dual filters. The attention term captures tokens with direct relational relevance, while the routing term acts as a structural safety net that rescues functionally specialized tokens (such as intermediate mathematical variables and programming keywords) that receive low direct query attention.

Full EKVA v2 Formulation: Incorporating the explicit attention sink preservation constraint ($\text{Sink}(x_t)$ for $t < 4$) and exponential recency decay ($\text{Recency}(x_t)$ with $\tau = 32$) establishes our peak performance: 57.41% on GSM8K ($+3.87$ pp over Attention-Only), 49.68% on HumanEval ($+3.18$ pp), and 90.48% on NIAH ($+5.89$ pp). Sink protection preserves initial prompt tokens that anchor the numerical stability of the softmax normalization denominator, while exponential recency ensures immediate contextual continuity for autoregressive token prediction. Together, all four terms in Eq. (10) are empirically necessary to maximize compression fidelity.

B. Expert Specialization Dynamics

During the offline calibration pass on $\mathcal{D}_{\text{calib}}$, tokens are categorized into four coarse syntactic classes $\mathcal{C}$ (numerics, keywords, syntax, and natural language) to trace the empirical dispatch distributions $P_{e,l}(c)$ across all layer-expert pairs. As visualized in the activation heatmap (Fig. 6), specific experts develop pronounced functional specializations ($\text{Spec}_{e,l} > 0.65$), particularly within intermediate transformer layers (e.g., layers 8 through 18). For instance, specific expert subsets preferentially activate for symbolic arithmetic operands, algebraic operators, and programmatic control keywords (`while`, `def`, `lambda`). When domain-critical tokens pass through the network during inference, they activate these specialized sub-circuits, resulting in elevated specialization scores that scale the routing niche metric $R(x_t)$ via Eq. (8). This mechanism guarantees that functionally indispensable tokens are retained in the shared Key-Value cache even when their query-key attention mass is temporarily muted.

VIII. OPERATIONAL SCOPE & LIMITATIONS

While EKVA v2 achieves consistent empirical gains and hardware speedups, we delineate the following operational scope and methodological limitations:

MoE Architectural Dependency: EKVA v2 fundamentally relies on capturing the discrete expert dispatch decisions generated by internal router gates during the prompt prefill phase. Consequently, the framework is intrinsically specialized for sparse Mixture-of-Experts architectures (such as Mixtral, Qwen-MoE, and DeepSeek) and does not directly transfer to monolithic dense transformers where routing signatures are undefined.

Shared Attention Topology Assumption: The mathematical formulation assumes standard open-weight MoE architectural conventions where self-attention is shared and dense across all sequence tokens, with expert routing restricted downstream to feed-forward blocks. While this topology characterizes virtually all contemporary foundation MoEs, future architectures that might introduce private, expert-specific self-attention projections would require decentralized per-expert cache eviction policies.

Calibration Overhead & Distribution Shifts: Although the offline calibration pass is lightweight (< 2 minutes over 40 sequences on a single GPU) and incurs zero inference-time runtime memory overhead, the computed expert profiles ($\bar{H}_{e,l}, \text{Route}_{e,l}, \text{Spec}_{e,l}$) assume that the calibration corpus adequately reflects the syntactic variety of downstream tasks. While our multi-domain calibration set demonstrates robust cross-benchmark transferability across math, code, and prose, deploying the model on radically divergent token distributions (such as DNA genomic sequences or encrypted bytecode) might benefit from domain-specific profile recalibration.

Decoding Regime: All benchmark evaluations were conducted under greedy decoding (`do_sample=False`) to provide deterministic, reproducible comparisons. While greedy decoding represents standard practice in mathematical and code benchmarks, assessing retention dynamics under stochastic nucleus sampling ($p = 0.9, T = 0.7$) represents a valuable future research inquiry.

IX. CONCLUSION & FUTURE WORK

We presented EKVA v2, a training-free token retention framework tailored specifically for Key-Value cache compression in sparse Mixture-of-Experts inference. By discovering and exploiting the empirical orthogonality between cross-layer routing signatures and self-attention mass ($\rho \approx 0.00$), EKVA v2 unifies routing signals with layer-specific expert niche profiles to preserve critical tokens in the shared KV cache without model fine-tuning or retraining. Across three open-weight MoE families (spanning 8, 60, and 64 experts), EKVA v2 delivers statistically significant 3.0 to 6.0 percentage point improvements ($p < 0.01$) over state-of-the-art baselines across mathematical reasoning, code generation, long-context retrieval, and language modeling under a 40% cache budget. Grounded in the Williams Roofline Model, our custom fused Triton

compaction kernel achieves a measured $2.02\times-2.05\times$ end-to-end decode step speedup on NVIDIA A100 GPUs.

Future work will investigate multi-batch Triton compaction kernels for continuous high-concurrency serving, extending routing-aware retention to vision-language MoEs, and exploring online adaptive recalibration of expert profiles during multi-turn conversations.

REFERENCES

- [1] T. Dao, “Flashattention-2: Faster attention with better parallelism and work partitioning,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [2] G. Xiao, Y. Tian, B. Chen, S. Han, and M. Lewis, “Efficient streaming language models with attention sinks,” in *International Conference on Learning Representations (ICLR)*, 2024.
- [3] Z. Zhang, Y. Sheng, T. Zhou, T. Chen, L. Zheng, R. Cai, Z. Song, Y. Tian, C. Ré, C. Barrett, Z. Wang, and B. Chen, “H2o: Heavy-hitter oracle for efficient generative inference of large language models,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 36, 2023.
- [4] Y. Li, Y. Huang, B. Yang, B. Venkitesh, S. Avestimehr, and M. Annavaram, “Snapkv: Llm knows what you are looking for before generation,” *arXiv preprint arXiv:2404.14469*, 2024.
- [5] Z. Feng *et al.*, “Ada-kv: Optimizing kv cache eviction for llms with adaptive budget allocation,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- [6] A. Chen *et al.*, “Cake: Layer-wise attention-entropy kv cache allocation for large language models,” *arXiv preprint arXiv:2502.04189*, 2025.
- [7] Z. Lin *et al.*, “Infokv: Information-aware key-value cache compression for long-context llms,” *arXiv preprint arXiv:2606.26875*, 2026.
- [8] A. Q. Jiang, A. Sablayrolles, A. Roux, A. Mensch, B. Savary, C. Bamford, D. S. Chaplot, D. d. I. Casas, E. B. Hanna, F. Bressand *et al.*, “Mixtral of experts,” *arXiv preprint arXiv:2401.04088*, 2024.
- [9] J. Bai *et al.*, “Qwen1.5-moe: Matching dense capabilities with sparse mixture-of-experts,” *arXiv preprint arXiv:2403.18731*, 2024.
- [10] D. Dai *et al.*, “Deepseek-moe: Towards ultimate expertise in mixture-of-experts language models,” *arXiv preprint arXiv:2401.06066*, 2024.
- [11] Y. Zhang, Z. Wang, J. Du, M. Zhang, and J. Zhao, “Pyramidkv: Dynamic kv cache compression based on pyramidal information funneling,” *arXiv preprint arXiv:2406.02069*, 2024.
- [12] R. Liu *et al.*, “Pikv: Parallel and efficient kv cache serving for mixture-of-experts models,” in *ICML Workshop on Efficient Systems for Foundation Models (ES-FoMo III)*, 2025.
- [13] H. Sun *et al.*, “Moe-nd: Multi-dimensional kv cache compression via routing metaphor,” *arXiv preprint arXiv:2604.17695*, 2026.
- [14] D. Balashov and A. Ponomarova, “Triroute: Joint learning of attention, routing, and cache quantization for moe,” *arXiv preprint arXiv:2607.06601*, 2026.
- [15] K. Cobbe, V. Kosaraju, M. Bavarian, M. Chen, H. Jun, L. Kaiser, M. Plappert, J. Tworek, J. Hilton, R. Nakano, C. Hesse, and J. Schulman, “Training verifiers to solve math word problems,” *arXiv preprint arXiv:2110.14168*, 2021.
- [16] M. Chen, J. Tworek, H. Jun, Q. Yuan, H. P. d. O. Pinto, J. Kaplan, H. Edwards, Y. Burda, N. Joseph, G. Brockman *et al.*, “Evaluating large language models trained on code,” *arXiv preprint arXiv:2107.03374*, 2021.
- [17] J. W. Rae, A. Potapenko, S. M. Jayakumar, and T. P. Lillicrap, “Compressive transformers for long-range sequence modelling,” *arXiv preprint arXiv:1911.05507*, 2019.
- [18] S. Williams, A. Waterman, and D. Patterson, “Roofline: An insightful visual performance model for multicore architectures,” *Communications of the ACM*, vol. 52, no. 4, pp. 65–76, 2009.